



PyTorch vs Tensorflow

Michael J. Williams

PHAS-ML September 21st 2020

Introduction

- This is based on my own experience and what I've heard/read from others
- There are other frameworks!
- This is Python centric



Shogun

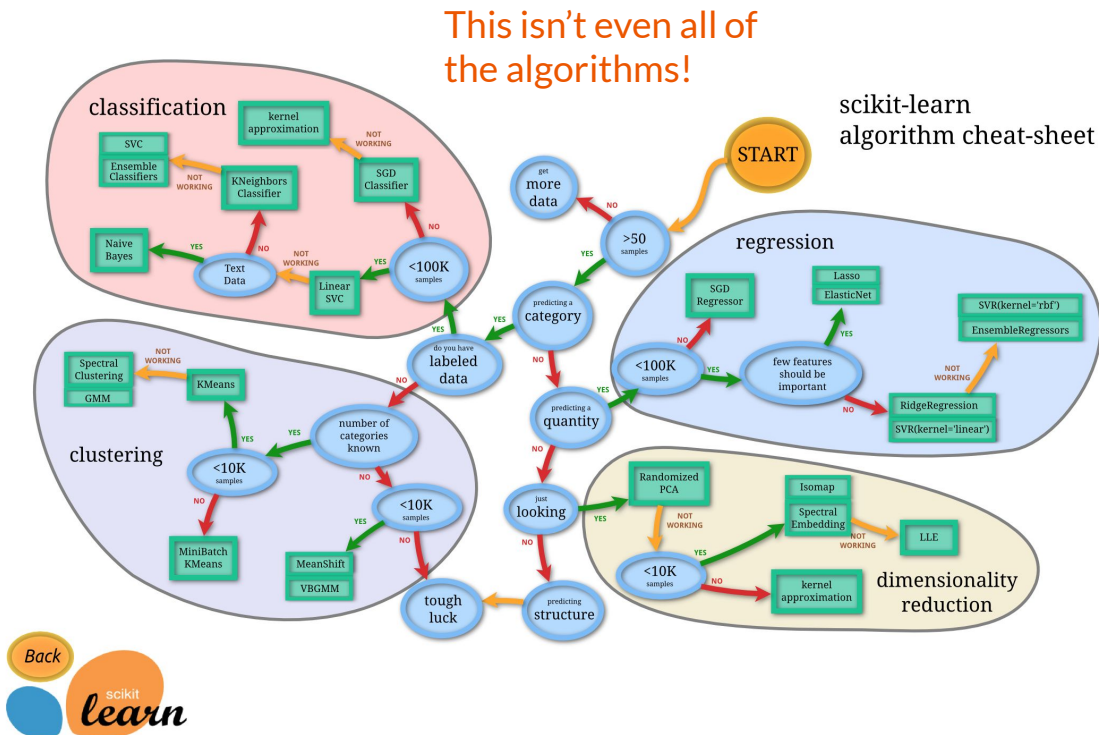


I'm new to ML, what should I use?

—

Scikit Learn

- Great starting point
- Simple and consistent
- Wide variety of algorithms
- Great documentation
- Lots of useful metrics for analysing results





When to try something else?

When you need something more complex:

- Deep neural networks
- GPU acceleration
- Higher dimensional data
- Parallelism and serialization

PyTorch and Tensorflow




Background



- Primarily developed by Facebook
- Initially released in Sept. 2016
- Originally based on Torch
- Some applications in industry
- Wide spread use in research



- Developed by Google Brain
- Initially released in Nov. 2015
- Used for all Google research
- Large application in industry



Comparison between unique mentions of PyTorch and Tensorflow in research journals.

Image credit: Horace He, "The State of Machine Learning Frameworks in 2019", The Gradient, 2019.




A note on TensorFlow 2.0

- Tensorflow 2.0 released in September 2019
- A lot changed, including the default behaviour of compute graphs which is now Eager just like PyTorch
- Many comparisons you'll read are between Tensorflow 1.x and PyTorch and most of the points no longer apply
- For anyone still using TF 1.x, upgrade if you can, it's a much better experience



Pros

-  Keras
- Tensorboard
- Deployment/production
- Static graph

Cons

- Documentation
- Debugging
- Static graph
- Installation



Pros

- [K Keras](#)
- [Tensorboard](#)
- Deployment/production
- Static graph

Cons

- [Documentation](#)
- [Debugging](#)
- Static graph
- Installation



PyTorch

Pros

- Documentation
- Debugging
- Dynamic graph
- Pythonic
 - Basic functions very similar to NumPy
- Installation

Cons

- Visualisation
- Deployment/production
- High level API requires 3rd party packages



PyTorch

Pros

- Documentation
- Debugging
- Dynamic graph
- Pythonic
 - Basic functions very similar to NumPy
- Installation

Cons

- Visualisation
- Deployment/production
- High level API requires 3rd party packages

Examples



A Simple CNN



```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

But what if I want to define my own Layer?

Or change how the forward pass works?



```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

A Simple CNN



```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Just tell Keras the number of neurons and the size of the filters and it sorts the rest!

Customising Keras models is often more difficult, you'll probably need to mix in base TF

Custom layers inherit from `nn.Module` and need a forward method



```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

You have to work these numbers out manually

You can do whatever you want here! Split inputs, have conditions, etc



Training



```
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))
```



```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
    running_loss = 0.0
```

Training



```
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])  
  
history = model.fit(train_images, train_labels, epochs=10,  
                    validation_data=(test_images, test_labels))
```

Single line for training

Manually compute the outputs and loss and then update the weights

'Callbacks' can be added and are called at the end of each epoch



```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
for epoch in range(2): # loop over the dataset multiple times  
  
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):  
        # get the inputs; data is a list of [inputs, labels]  
        inputs, labels = data  
  
        # zero the parameter gradients  
        optimizer.zero_grad()  
  
        # forward + backward + optimize  
        outputs = net(inputs)  
        loss = criterion(outputs, labels)  
        loss.backward()  
        optimizer.step()  
  
        # print statistics  
        running_loss += loss.item()  
        if i % 2000 == 1999: # print every 2000 mini-batches  
            print('[%d, %5d] loss: %.3f' %  
                  (epoch + 1, i + 1, running_loss / 2000))  
            running_loss = 0.0
```

No callbacks, but can easily define custom operations during training



Links to these examples

PyTorch

- https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

Keras

- <https://www.tensorflow.org/tutorials/images/cnn>

TensorFlow

- <https://www.tensorflow.org/tutorials/customization/basics> (three tutorials on customising datasets, layers and training)

Things to try:

- Saving model
 - Then saving only best model
- Learning rate scheduler
- Early stopping (patience)
- Plotting during training (this is harder in Keras)



More tutorials

Both have great tutorials that cover most of the basics:

 PyTorch

- <https://pytorch.org/tutorials/>

 Keras

- <https://www.tensorflow.org/tutorials>



Some interesting packages





- Lightning and Ignite (Higher level API, removes some boilerplate):
<https://github.com/PyTorchLightning/pytorch-lightning>,
<https://github.com/pytorch/ignite>
- BoTorch (Bayesian optimisation in Torch):
<https://botorch.org/>
- Pyro (Probabilistic programming):
<http://pyro.ai/>



- TF Probability:
<https://www.tensorflow.org/probability>
- TF Hub (use start-of-art pre-trained models): <https://www.tensorflow.org/hub>
- TF Quantum:
<https://www.tensorflow.org/quantum>



Conclusions

- Pytorch and Tensorflow are becoming more and more similar
- With Keras it's trivial to build and training a NN, however customisation can be more difficult since this requires use of more of the core TF functionality
- PyTorch doesn't have an equivalent high level API built in, but consequently customisation is often easier
- I'd say it comes down to personal preference and the application in question.
- For me:
 - Easy deployment:  TensorFlow
 - Flexibility and easy development in a research environment:  PyTorch

What do you think?

—